Handlers spend their time on judgement, not on reading and retyping.

1.85 million

complaints reached UK firms in the first half of 2026, each with an 8-week clock (FCA).

One window

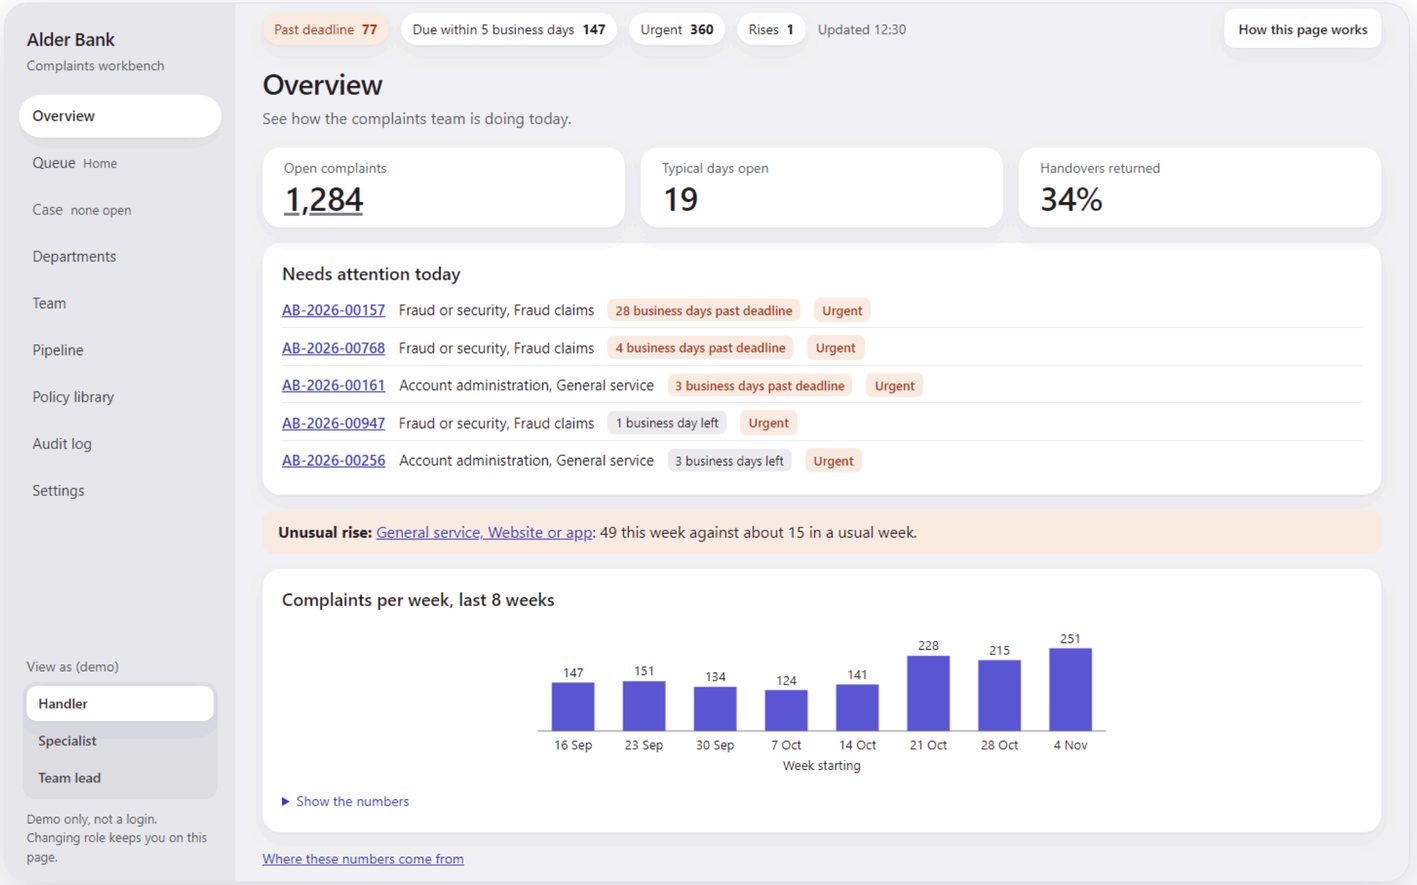
shows every deadline, what is due next and what needs this person, with one click to act.

90.5%

routed to the right team when the model is confident; the 11.3% it is unsure of go to a person.

8 clicks

from the queue to an approved, checked handover on the demo case (filter clicks not counted).

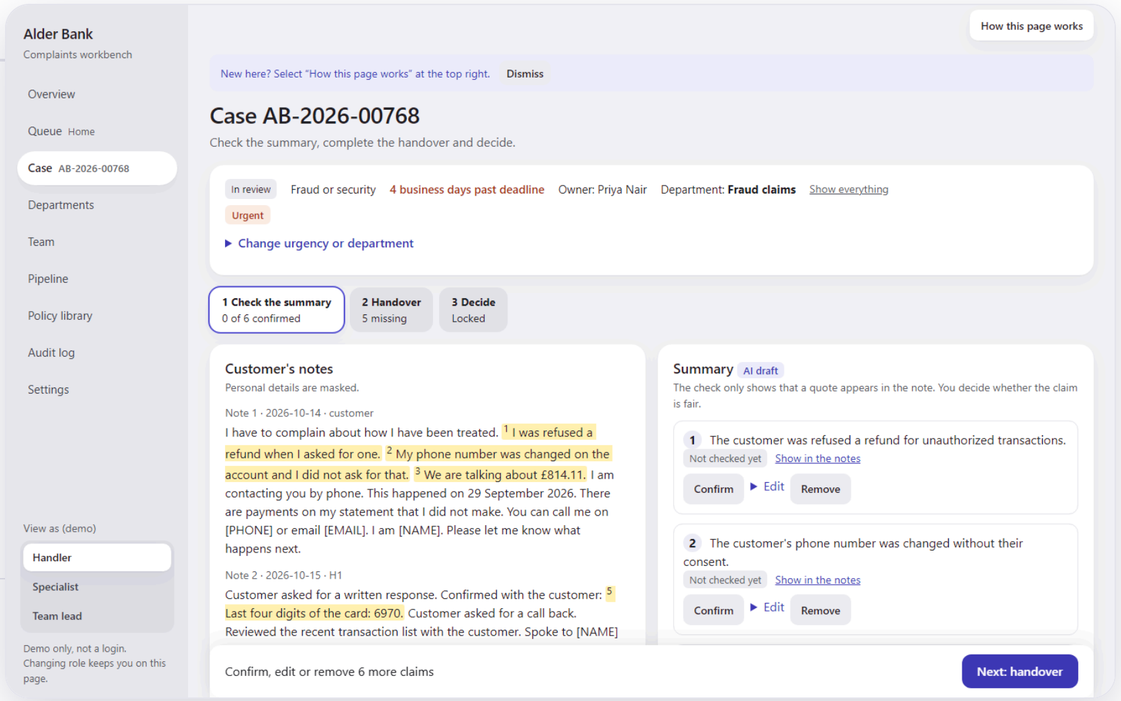


Value comes from time saved per case and fewer handovers sent back. Neither is measured on real work yet.

RISKS, CONTROLS AND HUMAN OVERSIGHT

The model drafts. A person decides, claim by claim.

Risk	Control	Measured on the test bench
An invented or wrong claim	Every claim needs a word-for-word quote from the notes; the handler confirms each one	**42 of 864 claims blocked**
Personal data reaches the model	Masked first; unmasking needs a reason and is logged	**21 of 22 test strings caught**
Sent to the wrong team	Low-confidence cases wait for a person	**11.3% go to a person**



Also: no bulk approve, urgent reasons shown in words, nothing leaves the machine (no network calls, strict content policy).

HOW THE WORK CHANGES

Handlers check instead of retype. Specialists get complete handovers.

Before	Reads all the notes	Searches policy by hand	Types the handover	Specialist finds gaps, sends back
With the tool	Masked draft with quotes	Handler checks each claim	Pack names missing fields	Accept, or return with a reason

Handler

Works a queue in deadline order with a preview and one-click actions.

New skill: checking AI output, not writing it.

Specialist

Sees missing fields before accepting.

Returns need a reason, so the handler learns why.

Team lead

Sees the whole flow and where cases come back.

Shares work out. No speed league table.

The "before" row is our description of a typical process, not a measurement.

WHAT WE MEASURED, AND WHAT IT DOES NOT PROVE

Every figure comes from one command, with its limit beside it.

Test bench	Result
Data handled	**1,600 cases: every page under 100 ms**
At 10 times the volume	**16,000 cases: Overview and Pipeline miss 100 ms**
Security	**Injection, tamper, cross-site and host checks pass; masking 21 of 22**
Ease of use	**9 pages x 3 roles fit one laptop screen**
Routing accuracy	**84.4% overall on held-out data**
Policy search	**13 of 20 right clause in the top five**
Tests	**TESTCOUNT automated tests, unit and browser**

Full results: docs/test-results.md. Run: uv run python tests/bench/bench.py --suites

Assumptions

All data is synthetic, shaped by public FCA, ombudsman and UK Finance figures; channel mix and time of day are our guesses. Settings lists each source.

Limits

A quote check proves the words exist, not that the summary is fair. Masking misses bare first names. The role switch is a demo, not a login.

Trade-offs

A small local model with drafts recorded in advance: reproducible and private, but not live. Classic models for routing: cheap and explainable.

AI TOOLS USED, AND HOW WELL THEY DID

AI did most of the typing. A person set direction; reviews caught what AI missed.

In the product

TF-IDF + logistic regression routing and type

SQLite FTS5, BM25 policy search

qwen3:4b-instruct drafts, run locally with Ollama and recorded

One stage in ten uses a language model.

In the build

Claude Code a controller (Claude Fable 5.2, then Claude Opus 5.5) ran builder and reviewer agents on cheaper tiers where the task allowed

Codex the case flow diagram and two second-opinion diagnoses

Local Qwen the recorded drafts

Assessment

Worked separate reviewers caught masking leaks, an audit gap, a page that wrote on read and a self-flattering accuracy (0.894 fell to 0.844).

Did not an agent misread "too busy" as "too many charts"; the owner corrected it. Codex could not commit from its sandbox.

Plan and spec → Builder agent: tests, then code → Read-only reviewer → Fix rounds → Owner looks and corrects

Build log: docs/loop-log.md. Code and instructions: README.md.