ankura

Research first. Every decision has a note.

Before any code, the research went into an Obsidian vault: complaint rules, AI risks, UX evidence, each fact checked at its source. Because every note links to the evidence behind it, the AI agents could follow a decision straight to its source instead of re-reading everything.



113	916	39	9
linked notes	links between them	decisions	research notes

01	02	03	04	05
Rules and evidence	Checked at source	One note per decision	Spec and 13-task plan	Build, test, review

Not an Ankura product · synthetic data only

ankura

Judgement, not retyping.

The workbench routes each complaint, drafts a summary that quotes the customer's own words, and builds the handover. People keep every decision.

1.85M
complaints to UK firms in H1 2026, each on an 8-week clock (FCA)

90.5%
routed to the right team when the model is sure

8
clicks from the queue to an approved, checked handover

11.3%
of cases the model is unsure of go straight to a person

01
Handler
Reads every note, retypes the handover.
→ Checks an AI summary claim by claim against its quotes.

02
Specialist
Finds gaps, sends cases back.
→ Sees missing fields before accepting; a return needs a reason.

03
Team lead
Chases deadlines case by case.
→ Sees the flow, the rises and the load, with no speed league table.

Value: time per case and fewer handovers sent back. Not yet measured on real work; "today" is our description.

The model drafts.
A person decides.

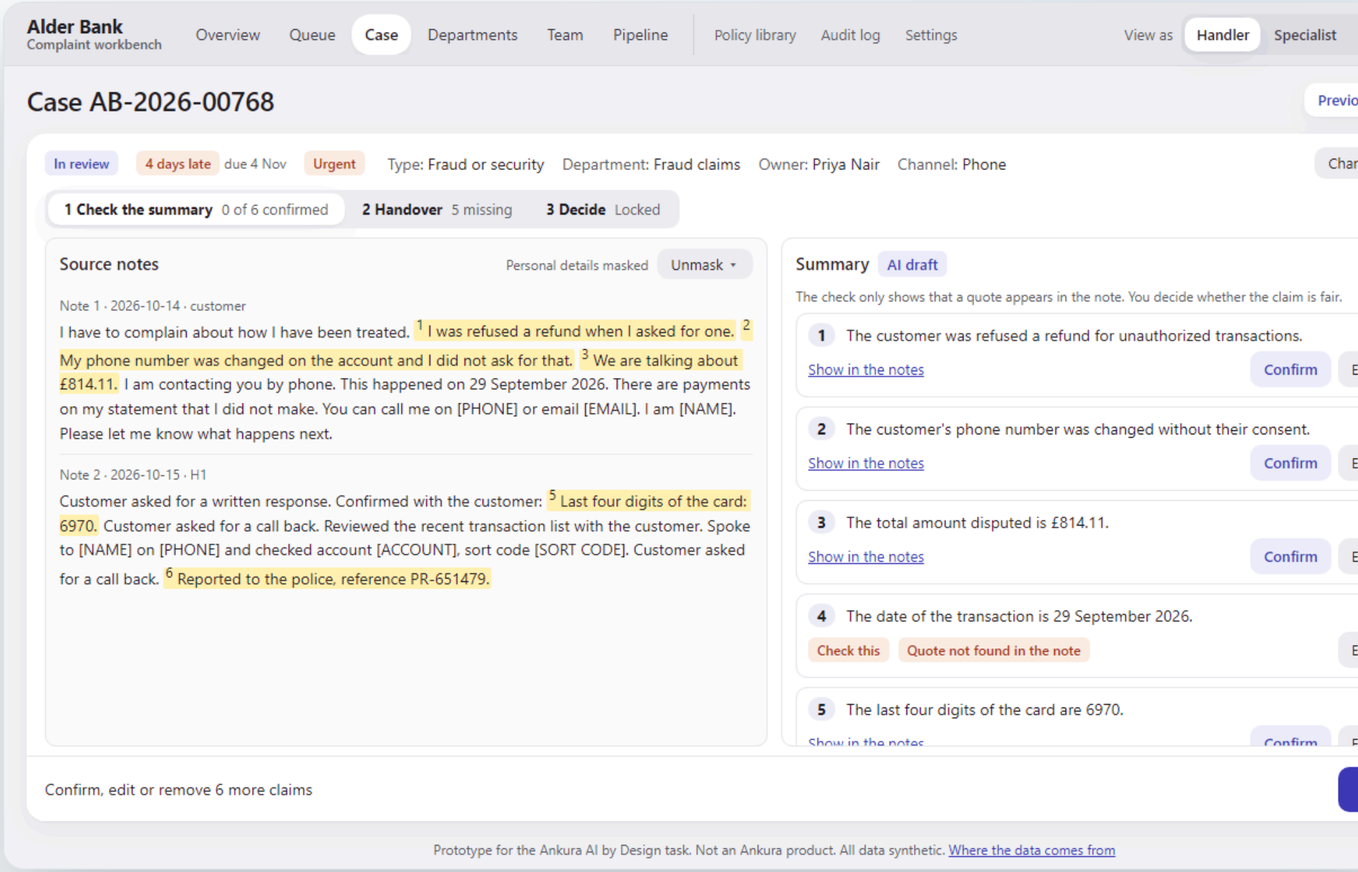
INVENTED CLAIMS — 42 of 864 blocked
Every claim must quote the notes word for word

PERSONAL DATA — 21 of 22 caught
Masked before the model; unmasking is logged

WRONG TEAM — 11.3% to a person
Unsure cases wait for a person

HIDDEN INSTRUCTIONS — 0 of 1 changed
Notes are fenced as data, output checked

SILENT EDITS — detected, not prevented
Every action in a hash-chained log



Claim by claim
no bulk approve

0
network calls in use

01

Assumptions

All data is synthetic, shaped by public FCA, ombudsman and UK Finance figures. Channel mix and time of day are our estimates; Settings lists every source.

02

Limits

A quote check proves the words exist, not that the summary is fair. Policy search finds the right clause 13 times in 20. The role switch is a demo, not a login.

03

Trade-offs

A small local model with drafts recorded in advance: private and repeatable, not live. Classic models for routing: cheap, fast and explainable.

1,600

synthetic cases, from a fixed seed

13 of 20

right policy clause in the top five

12 s

per draft on a laptop, recorded once

One stage in ten uses a language model

TF-IDF and logistic regression · BM25 search · a small Qwen model, run locally

Routing and search use classic methods; the model only drafts, and its drafts are checked.

A controller and a team of agents

Claude Code · Codex · a local Qwen model

Builder and reviewer agents worked from a written plan and a ledger, four builders in parallel at the peak, each on the cheapest model that could do the job.

Fast, but only as good as its checks

Worked. Separate reviewers caught masking leaks, an audit gap and an accuracy figure that flattered itself (0.894 fell to 0.844).

Did not. Codex could not commit from its sandbox; an outage stopped every agent at once.

Figures and misses: docs/test-results.md · Build log: docs/loop-log.md

Built and tested like a product.

527

automated tests, unit and browser, all passing

63

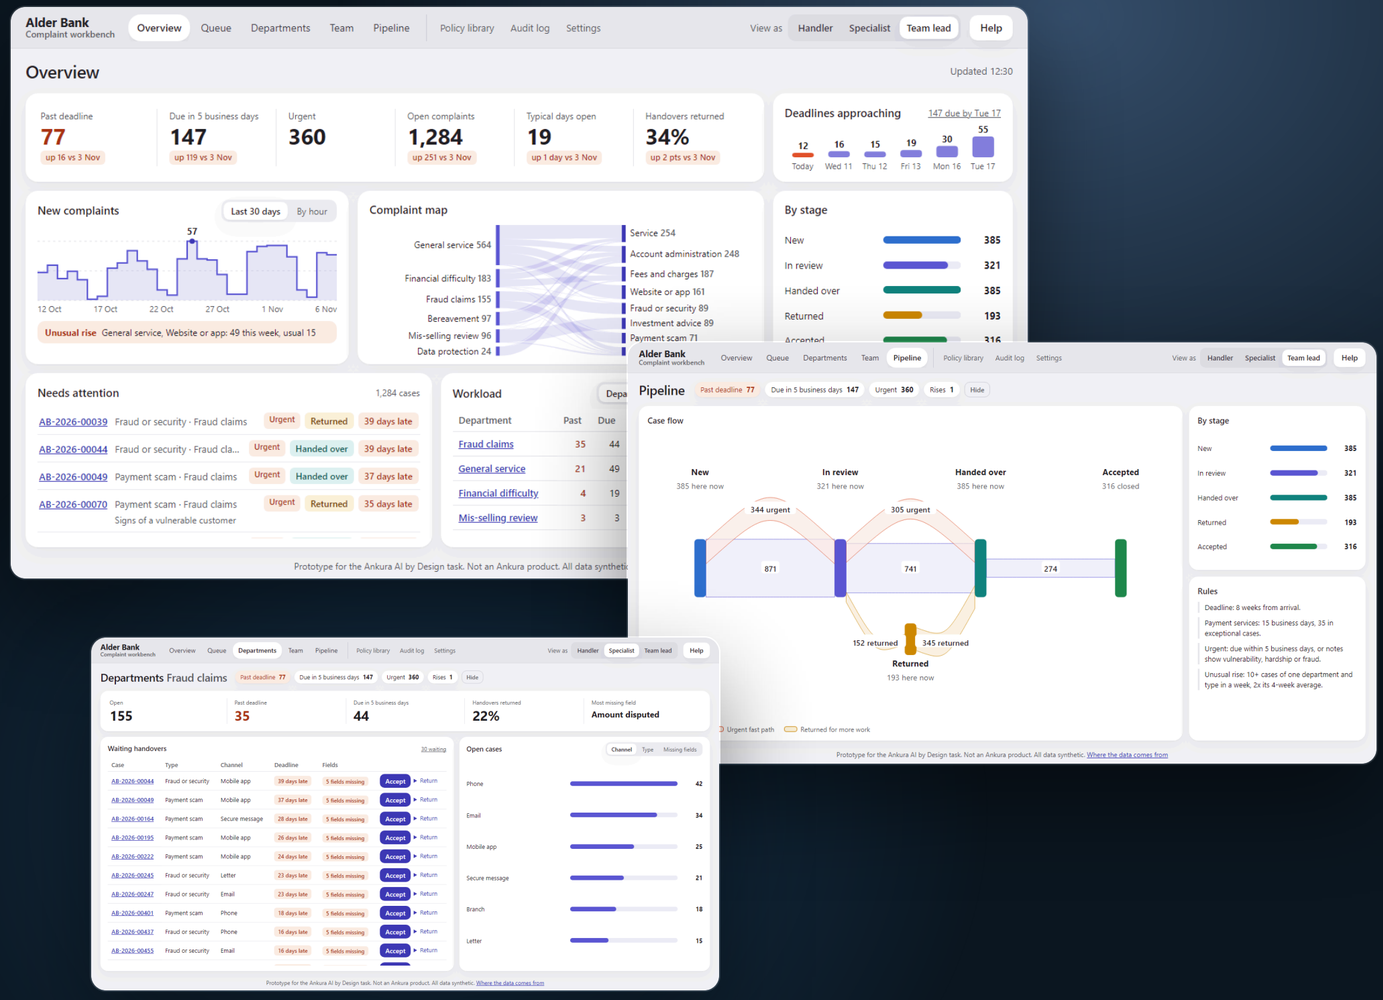
browser checks that 9 pages fit one screen for 3 roles

<100 ms

every page at 1,600 cases; the Overview takes 345 ms at 16,000

0

outside calls while in use; strict content policy, no build step



Run it: uv sync, then uv run workbench · Bench: uv run python tests/bench/bench.py --suites